

Assignment 2: HTTP Server

CSE 130: Principles of Computer Systems Design

Due: May 5, 2023 at 11:59 PM

Goals This assignment will provide you with experience building a system that uses client-server/strong modularity. Our learning objectives are: (1) practice implementing a client-server system, (2) practice the advantages of powerful abstractions, (3) practice implementing a large system that solves a large problem, (4) review string parsing, and (5) review memory management.

This assignment is more complex than Assignments 0 and 1. This doc is long out of necessity: we have a lot of things to explain! We encourage you to first, read and annotate this document early (in the first few days that the assignment is released), then, design your server with a “box-and-arrows” diagram (within the first week that the assignment is released), and only then start writing code.

Assignment Details

You will be building an HTTP server for this assignment. Your server should execute “forever” without crashing (i.e., it should run until a user types `CTRL-C` on the terminal). Your Server will create, listen, and accept connections from clients that arrive on a *port*. Your server’s task is to process the bytes sent by clients using the HTTP protocol. A key task in this assignment is building a server that is resilient to malformed (or even malicious) clients: *no matter what a client sends your server, your server should not crash!*

In this section, we next walk through (1) how to run your server, (2) what your server’s execution should perform, and (3) provide a few examples to help you understand.

Running your server

Your server should take a single-command line argument, an `int`, named *port*. In other words, your server should be started by specifying the following command from the `asn2` directory of your repository:

```
./httpserver <port>
```

Your Server’s Execution

This section provides an overview of the tasks that your server should perform. First, your server should create structures needed to listen for incoming connections. Then, it should repeatedly accept a new client connection, process that connection, and close that connection. Note, this should occur one connection at a time (i.e., it should accept and process its first connection before it begins processing the second connection). We outline the requirements for each of these steps below:

Listening for Connections

First, your server should create a *socket*, bind the socket to a `port`, and make it listen for incoming connections. We will provide you with a struct definition, `Listen_Socket`, for this socket. Additionally, we will provide a function, with `int listener_init(*Listen_Socket, int port)`, that initializes a `Listen_Struct` to listen to a given port¹. If the `port` passed to `httpserver` is invalid (not an integer between 1 and 65535) or if `httpserver` cannot bind to the provided `port`, then your `httpserver` should produce the message “Invalid Port\n” to `stderr` and exit with a return code of 1. Check out the return value of `listen_init`; it might help you with this task...

¹Note: We encourage you to use our library functions, but you don’t *have* to do so if you’d rather build everything yourself.

Accepting Connections

Your server should repeatedly accept connections made by clients to the `port`. Linux represents a connection with an external entity (in this case, a client) using the `socket` abstraction. We will provide you with a library function, `int listener_accept(*Listener_Socket sock)`, that blocks execution until a new incoming connection is made to the socket, `sock`, and then returns a new socket for the new connection.

After accepting a connection, your server should process any valid HTTP commands that the client sends on that connection. To do this, your server will need to read bytes sent by the client and send bytes to the client. The `socket`, created when your server accepts the new connection, makes this task possible. You should think of a socket as equivalent to a file descriptor: It is an integer. Your server can read bytes in the order that they were written from the other side of the socket (i.e., the client) by calling `read` with the socket as the `fd` argument. Your server can write bytes in the order that they were written from the other side of the socket (i.e., the server) by calling `write` with the socket as the `fd` argument.

Processing Connections

Your server should process a simplified subset of the HTTP Protocol². In particular, your server should support two types of HTTP operations: `GET` and `PUT`. Below, we describe the HTTP request format that your server should be able to handle, the HTTP response format that your server should send to clients as a response, and then describe the processing for `GET` and `PUT` commands.

Requests `GET` and `PUT` requests share the same request format (i.e., a valid `GET` request and a valid `PUT` request contain the same fields). Below we show the format, identifying those that are required and those that are Optional. Then, we describe each field in detail. Note, each field is separated, without any additional whitespace or characters, by the sequence `\r\n`. This sequence is the sequence of the single character ‘`\r`’ followed by the character ‘`\n`’, **NOT** the sequence of four characters ‘`\`’, followed by ‘`r`’, followed by ‘`\`’, followed by ‘`n`’:

```
Request-Line\r\n    ; [Required]
(Header-Field\r\n)* ; [Optional, repeated]
\r\n
Message-Body      ; [Optional]
```

A valid request’s characters up to and including the “double” `\r\n` will not exceed 2048 characters (this includes the length of the `Request-Line`, all `Header-Fields`, and all `\r\n` delimiters. Below, we describe each of the fields of a request in detail:

- **Request-Line.** Every valid request includes exactly one request line with the following format:

`Method URI Version`

where each field, (`Method`, `URI`, and `Version`) are case-sensitive sequences of characters. A valid request line will contain at most 2048 characters. We limit valid requests in the following ways:

- A valid `Method` contains at most eight (8) characters from the character range `[a-zA-Z]`. Your server only needs to implement (i.e., perform the semantics) of `GET` and `PUT`. (The distinction between “treats as valid” and “implements” will be clear after the Responses section below).
- A valid `URI` starts with the character ‘`/`’, includes at least 2 characters and at most 64 characters (including the ‘`/`’), and with the exception of the leading ‘`/`’, only includes characters from the character set `[a-zA-Z0-9.-]` (this character set includes 64 total valid characters). Each `URI`, `/path`, matches the file `path` within the directory in which the `httpserver` is running. For example, the `URI /foo.txt` matches the file `foo.txt` from the folder in which the `httpserver` is running. **Note: your server does not need to handle path’s that include directories.** Your server should only perform the semantics of `GET` and `PUT` on requests that include a valid `URI`.
- A valid `Version` has the format `HTTP/#.#`, where each `#` is a single digit number. Your `httpserver` should only implement version 1.1, so it should only perform the semantics of `GET` and `PUT` requests that include a `version` equal to `HTTP/1.1`.

²The curious reader can find the whole HTTP Protocol at Request For Comments: 2616 (abbreviated RFC 2616). We’ve made a number of simplifications, though, so follow our instructions rather than the RFC.

Status-Code	Status-Phrase	Message-Body	Usage
200	OK	OK\n	When a method is Successful
201	Created	Created\n	When a URI's file is created
400	Bad Request	Bad Request\n	When a request is ill-formatted
403	Forbidden	Forbidden\n	When a URI's file is not accessible
404	Not Found	Not Found\n	When the URI's file does not exist
500	Internal Server Error	Internal Server Error\n	When an unexpected issue prevents processing
501	Not Implemented	Not Implemented\n	When a request includes an unimplemented Method
505	Version Not Supported	Version Not Supported\n	When a request includes an unsupported version

Table 1: List of HTTP status that your `httpserver` should support.

- **Header-Field.** Valid requests include zero (0) or more `header-fields` after `request-line`. A `header-field` is a key-value pair with the format:

`key: value\r\n`

The `key` ends with the first instance of a `'.'` character. A valid request's header-field keys will be at least 1 character, at most 128 characters, and only contain characters from the character set `[a-zA-Z0-9.-]`. A valid request's header-field values will contain at most 128 characters and only contain characters from the set of printable ASCII characters (i.e., a valid value will not contain any ASCII "Device Control" characters nor any other binary data).

Valid requests separate each `header-field` using the sequence `\r\n`, and will terminate the list of `header-fields` with a blank header terminating in `\r\n`. (Essentially, regardless of how many `header-fields` a request contains, the list will terminate with the sequence `\r\n\r\n`).

- **Message-Body.** Valid `PUT` requests must include a message-body; valid `GET` will not include a message-body. Valid requests that include a `Message-Body` will also include a header, with a `key` of `Content-Length`, whose value will indicate the number of bytes in the `Message-Body`.

As a summary, here is an example of a valid `PUT` request to the URI, `foo.txt`:

```
PUT /foo.txt HTTP/1.1\r\nContent-Length: 12\r\n\r\nHello world!
```

Responses Your `httpserver` must produce a response for each request, regardless of whether the request is valid or not. Your response must follow the grammar:

```

Status-Line\r\n      ; [Required]
(Header-Field\r\n)*  ; [Optional, repeated]
\r\n
Message-Body         ; [Optional]
```

- **Status-Line** The status line indicates the type of response to the request. It consists of three fields:
`HTTP-Version Status-Code Status-Phrase`
`httpserver` must always produce the `HTTP-Version` string, `HTTP/1.1`, regardless of the `HTTP-Version` provided in the request. A response with a `Status-Code` in the 200s indicates a successful response, in the 400s indicates an erroneous response, and in the 500s indicates an internal server error. Table 1 lists the `status-codes` that `httpserver` needs to produce, and their associated `status-phrase` that your server should produce in the response's status line, the message body that your server should produce as a part of the response, and when you should use each code.
- **Header-Field.** The `status-line` should be followed by zero (0) or more `header-fields`. A response's `header-fields` have the same format as the request header fields, namely:

key: value

Your server only needs to produce one header: each response should include a **Content-Length** header whose value is equal to the size of their **message-body**. Your server should separate each **header-field** using the sequence `\r\n` and terminate the list of **header-fields** with a blank header terminating in `\r\n`. (Essentially, the list of **header-fields** ends with `\r\n\r\n`).

- **Message-Body.** **httpserver** must produce a **message-body** with each response, whose size, in bytes, is equal to the value identified in the response's **Content-Length** header. Your server should produce the **Message-Body** indicated in Table 1 for each request, *except for valid get requests*. We describe the correct **message-body** for valid GET requests in the section on Methods.

Methods

You must implement two HTTP methods, GET and PUT.

- **GET** A GET request indicates that the client would like to receive the content of the file identified by the URI. If a request is valid and specifies a URI that is resident in the directory in which **httpserver** is executing, then **httpserver** should produce a response that...

1. has a **status-code** of 200
2. has a **message-body** that includes the current state of the file pointed to by URI, and
3. has a **Content-Length** that indicates the number of bytes in the file.

For all other requests (include those that are valid but where the URI indicates a non-existent file), your server should produce a **status-code**, **message-body**, and **Content-Length** based upon Table 1.

- **PUT** A PUT request indicates that the client would like to update/replace the content of the file identified by the URI. If a valid PUT request's URI points to a file that does not yet exist, **httpserver** should...

1. create the file
2. set the file's contents equal to the **message-body** in the request
3. produce a response with a **status-code**, **message-body**, and **Content-length** for the status-code 201 based upon Table 1.

If a valid PUT request's URI points to a file that does already exist, **httpserver** should...

1. replace the file's contents with the **message-body** in the request
2. produce a response with a **status-code**, **message-body**, and **Content-length** for the status-code 200 based upon Table 1.

For all other requests, your server should produce a **status-code**, **message-body**, and **Content-Length** based upon Table 1.

Closing Connections

After your server finishes processing a request for a connection, your **httpserver** should close the connection. You can simply call **close** on a socket like you would any other file descriptor. Note: each connection will contain at most one valid request; if it contains extra bytes after a valid request, your server should ignore those bytes (no need to send an additional response!). Also: before closing each connection, your server should read all of the bytes that were sent by the client, regardless of whether the client sent a valid request in their connection.

Notes

In addition to supporting the methods listed above, your project should meet the following functionality and limitations:

- **httpserver** should produce responses with the appropriate status code (see **Usage** in Table 1).

- `httpserver` should consider any connection in which it waits for 5 seconds for client input as issuing an invalid request. This is a timeout. Our helper functions are designed to help you with this requirement: `listener_accept` creates sockets for each connection that will return an `errno` of either `EAGAIN` or `EWOULDBLOCK` when a call our helper functions `read_until`, `read_all`, or `pass_bytes`, or to the system call `read`, experiences this timeout.
- `httpserver` should not have any memory leaks.
- `httpserver` will need to handle connections that do not send data (this is very important for our testing scripts)!
- `httpserver` should not leak any file descriptors.
- `httpserver` must be reasonably space efficient: it should use less than 10 MB of memory regardless of input.
- `httpserver` must be reasonably time efficient.
- `httpserver` should never crash (e.g., it should never `segfault`).
- `httpserver` must be written using the ‘C’ programming language (*not C++!*).
- `httpserver` cannot use the following functions from the ‘C’ `stdio.h` library: `fwrite`, `fread`, variants of `put` (i.e., `fputc`, `putc`, `putc_unlocked`, `putchar`, `putchar_unlocked`, and `putw`), and `get` (i.e., `fgetc`, `getc`, `getc_unlocked`, `getchar`, `getchar_unlocked`, `getline`, and `getw`).
- `httpserver` cannot use functions, like `system(3)`, that execute external programs.

Examples

In this section, we describe a few example requests and the correct responses for those requests. For each example, assume that your `httpserver` is started in a directory that contains the file, `foo.txt`, with content, “Hello World, I am foo”, and a file, `bar.txt` containing the content, “Hello World, I am bar”.

1. **Ex. 1** The client makes a GET request by sending:

```
GET /foo.txt HTTP/1.1\r\n\r\n
```

The server does not change any files. It responds to the client by sending the client the following response:

```
HTTP/1.1 200 OK\r\nContent-Length: 21\r\n\r\nHello World, I am foo
```

2. **Ex. 2** The client makes a PUT by sending:

```
PUT /foo.txt HTTP/1.1\r\nContent-Length: 21\r\n\r\nHello foo, I am World
```

The server replaces the contents of `foo.txt` with “Hello foo, I am World” and respond by sending the client the following content

```
HTTP/1.1 200 OK\r\nContent-Length: 3\r\n\r\nOK\r\n
```

3. **Ex. 3** The client makes a PUT request by sending:

```
PUT /new.txt HTTP/1.1\r\nContent-Length: 14\r\n\r\nHello\nI am new
```

The server creates a new file in its directory, named `new.txt`, with the contents “Hello\nI am new”. It then responds to the client with the following:

```
HTTP/1.1 201 Created\r\nContent-Length: 8\r\n\r\nCreated\r\n
```

4. **Ex. 4** The client makes a GET request by sending:

```
GET /not.txt HTTP/1.1\r\n\r\n
```

Since `not.txt` does not exist in the server's directory, the server responds with the following:

```
HTTP/1.1 404 Not Found\r\nContent-Length: 10\r\n\r\nNot Found\r\n
```

5. **Ex. 5** The client makes an invalid request by sending:

```
GET /foo.txt HTTP/1.10\r\nhello*world: value\r\n\r\n
```

This request is invalid because (1) `HTTP/1.10` is an invalid `Version` since 10 is not a single digit number, AND (2) The `header-field` is invalid since the key `hello*world` contains the `*` character. So, the server will respond by sending the client the following content:

```
HTTP/1.1 400 Bad Request\r\nContent-Length: 12\r\n\r\nBad Request\r\n
```

Rubric

We will use the following rubric for this assignment:

Category	Point Value
Makefile	10
Clang-Format	5
Files	5
Functionality	80
Total	100

Makefile Your repository includes a Makefile with the rules `all` and `httpserver`, which produce the `httpserver` binary, and the rule `clean`, which removes all `.o` and binary files. Additionally, your Makefile should use clang (i.e., it should set `CC=clang`), and should use the `-Wall`, `-Wextra`, `-Werror`, and `-pedantic` flags (i.e., it should set `CFLAGS=-Wall -Wextra -Werror -pedantic`).

Clang-Format All `.c` and `.h` files in your repository are formatted in accordance with the `.clang-format` file included in your repository.

Files The following files are included in your repository: `httpserver.c`, `Makefile`, and `README.md`. Your repository should not include binary files nor any object files (i.e., `.o` files), except for the file named `asn2_helper_funcs.a` (see Resources below). To make it easier for you to maintain tests, you can also include binary files in any directory whose name starts with the phrase `test`.

Functionality Your `httpserver` program performs the functionality described in Assignment Details.

How to submit

Submit a 40-character commit ID hash on the Canvas assignment to identify the commit that you want us to grade. We will grade the last hash that you submit to the assignment on Canvas and will use the timestamp of your last upload to determine grace days. For example, if you post a commit hash 36 hours after the deadline, we will subtract 2 grace days from your total.

Resources

Here are some resources to help you:

Testing

We provided you with two resources to test your own code:

1. An autograder, which is run each time you push code to GitLab, will show you the points that you will receive for your Makefile, Clang-Format, and Files.
2. A set of test scripts in the resources repository to check your functionality. You can use the tests to see if your functionality is correct by running them on your Ubuntu 22.04 virtual machine. We provided you with a subset of the tests that we will run, but, I bet you can figure the other ones out by adapting what we have given you :-)
3. A set of helper functions in the resources repository that will help you perform basic tasks. Table 2 contains a description of these functions. We released two things, `asgn2_helper_funcs.h`, a header file that declares each of the functions, and `asgn2_helper_funcs.a`, a Linux archive that contains each of the function definitions in a binary format. You can treat `asgn2_helper_funcs.a` as if it were a `.o` file in your makefile. For example, to build an executable, `httpserver`, using the files `httpserver.o` and `asgn2_helper_funcs.a`, you might execute:

```
clang -o httpserver httpserver.c asgn2_helper_funcs.a
```

Hints

Here are some hints to help you get going:

- You will likely need to look up how some system calls (e.g., `read`) and library functions (e.g., `warn`) work. You can always Google them, but you might also find the man pages useful (e.g., try typing `man 2 read` on a terminal).
- You will have to do a fair amount of string parsing for this assignment. There are many options on how to parse a string; we suggest that you look into using regular expressions. ‘C’ provides a `regex` library: it is a bit clunky, but it is significantly more powerful than using alternatives (such as, e.g., `strtok_r` or `sscanf`).
- There are a few ways to test `httpserver`. Below, we assume that you started `httpserver` on port 1234 by executing the command `./httpserver 1234`. We also assume that you are using your client on the same machine upon which the server is currently executing:
 - You can use an HTTP Client, such as Firefox, Google Chrome, or Safari. These can be fun ways to see your server in action! But, they are extraordinarily robust: our experience is that such web clients will happily accept responses that are formatted incorrectly.
 - RECOMMENDED: You can test with a command-line web browser, such as `curl`. `curl` can produce both `GET` and `PUT` commands. For example, to execute a `GET` of the file `foo.txt` on `httpserver` and place the output into the file `download.txt`, you execute:

```
curl http://localhost:1234/foo.txt -o download.txt
```

Use the following `curl` command to execute a `PUT` request that puts the file, `foo.txt`, into the location `new.txt`:

```
curl http://localhost:1234/new.txt -T foo.txt
```

Note: `curl` has a few idiosyncrasies that you might accidentally rely upon in your assignment. In effect, we’ve seen students in the past struggle on this assignment because of assumptions that they baked into their server which are based upon supporting `curl` as a client. We have two suggestions: (1) In addition to using `curl`, use `cse130_nc` (see below) to have even more control over what bytes are sent to your server and (2) In addition to using the basic `curl` commands shown above, also try clearing the `Expect` header by adding the following to the command-line `-H "Expect:"`

Signature	Description
<code>int listener_init(*Listener_Socket sock, int port)</code>	Initializes <code>sock</code> to be bound to <code>localhost</code> and listen on <code>port</code> . Note <code>sock</code> must already be allocated (i.e., it should either already be allocated on the stack or heap). Returns 0 if it was successful and 1 if it failed (either because <code>port</code> was in an invalid range, or because it could not bind to <code>port</code>).
<code>int listener_accept(*Listener_Socket sock)</code>	Blocks until a new client connection. Returns a <code>socket</code> for the new client connection; initializes a 5-second timeout on <code>socket</code> .
<code>int read_until(int fd, char *buf, int size, char *str)</code>	Reads bytes, placing them into <code>buf</code> , from the socket or file descriptor <code>fd</code> , until any of the following are true: (1) it has read exactly <code>size</code> bytes, (2) there are no more bytes to read from <code>fd</code> (i.e., <code>read</code> returned 0), or (3) the string, <code>str</code> is contained in <code>buf</code> . Returns the number of bytes read, or -1 if there was an error. If an error occurred, this function sets <code>errno</code> to be equal to the error that was encountered (see <code>errno.h</code>).
<code>int write_all(int fd, char *buf, int size)</code>	Writes exactly <code>size</code> bytes from <code>buf</code> into the socket or file descriptor <code>fd</code> , unless it encounters an error. Returns 0 if successful, or -1 if there was an error. If an error occurred, this function sets <code>errno</code> to be equal to the error that was encountered (see <code>errno.h</code>).
<code>int pass_bytes(int infd, int outfd, int size)</code>	Writes exactly <code>size</code> bytes from the socket or file descriptor <code>infd</code> into the socket or file descriptor <code>outfd</code> . Returns 0 if successful, -1 if there was an error reading from <code>infd</code> , and -2 if there was an error writing to <code>outfd</code> . If an error occurred, this function sets <code>errno</code> to be equal to the error that was encountered (see <code>errno.h</code>).

Table 2: Helper functions provided through resources.

- **RECOMMENDED:** You will find a file, `cse130_nc`, in the resources repository. This is our port of a standard Linux utility, `nc` or “netcat”. To connect to your server, execute `./cse130_nc localhost 1234`. Then, you can type in the text that you wish to send to your server. You can also automate this approach by piping data to `cse130_nc`. For example, to send the PUT command listed above to your server, execute:

```
printf "PUT /foo.txt HTTP/1.1\r\nContent-Length: 12\r\n\r\nHello World!"
| ./cse130_nc localhost 1234
```

- If you try to start your server immediately after killing a previous instance of it, you will likely see the following error:

```
httpserver: bind error: Address already in use
```

In this case, just restart the server with a different port number. The issue is that the operating system must ensure unique ports are used across the entire system; it often waits to gracefully close ports even after the process that was using them terminates.

Getting Started and Design Tips

This is a pretty large project and there are definitely multiple ways to get started. We describe an approach below based upon our experience working with students in the past, but you do not *have* to follow this section (that is why it’s under Hints!)

Design Approach

You should design your server before you start writing any code. Your design should articulate exact function prototypes and struct definitions that you plan to use. You may find it helpful to visualize how your server will function by drawing a picture similar to the “box-and-arrows” pictures that we’ve been drawing in class. By thoroughly designing your system, you will find that you can avoid many potential bugs, thereby saving you hours!

Keep in mind that your design will probably change as you build your server; sometimes, the act of trying to implement something illuminates how bad of an idea it was. If this happens, we encourage you to “go back to the drawing board” and re-design the server. If you find yourself implementing something that you haven’t “designed”, then you’re probably doing it wrong :-)

Step 1: Design your Scaffolding Start thinking about high-level “scaffolding”—essentially, what will your `main` function do? At a high-level, `main` will first need to parse command-line arguments and create a listener socket for the port. Then, `main` will need to repeatedly: (1) accept a new connection, (2) perform the logic (i.e., process) that connection, and (3) close that connection (“repeatedly”—sounds like a loop!). You should probably use “top-down” modularization. Make each of these tasks (getting command line arguments, accepting connections, processing connections, and closing connections) a separate function. Don’t worry about *how* those functions work, just treat them as a “black box” that performs the task you expect of them!

Step 2: Design your Connection Processing There are many ways to approach connection processing; we lead you through one approach.

1. **Read enough bytes to ensure that you will have read the entire Request-Line from the request.** Since your `httpserver` needs to be efficient, you will need to read multiple bytes at a time and place them into a buffer.

There is one important difference between a socket and a file descriptor that matters for this step. Recall that `read` returns 0 when a program reaches the end of a file. In contrast, a client will not always tell your server that it is done sending data. So, you might find yourself in a case where there is no new data to `read` from the socket, but the socket nonetheless does not return 0. If you use our helper functions in your assignment, you will eventually observe a timeout during `read`. But, you should not wait for that timeout before processing a correct request.

Think: How should you change how you read bytes from a connection to tell if you have gotten the full Request-Line? (A hint: Does the `read_until` helper function help you?)

2. **Identify each of the fields within the buffer that you read.** You'll want to next identify each of the fields that are contained in the buffer that you just read. *Think: What fields could be present in the buffer?*

There are many ways to parse the fields. For example, you could first separate the buffer into each high-level field (**Request-Line**, **Header-Fields**, **Message-Body**, (Can each of these be in your buffer?)) and then parse out sub-fields (e.g., **Method** within **Request-Line**). Or, you could parse each of the individual fields one by one from the beginning (i.e., start by parsing **Method** instead of **Request-Line**). Both of these approaches have merits; it is a personal preference what you choose. No matter what you choose, we suggest pulling out the 'C' regular expression library. It is clunky and awkward, but it is the best mechanism for the task at hand.

One important thing to decide upon is where you are going to store each of these fields. We suggest putting them into a struct. No matter where you put the fields, you should ensure that the lifetime of the data (the time that it is allocated, either through `malloc` or by virtue of being on the stack) covers the entire time that the fields are used. In other words: if you use a field in a function, `foo`, then the field must be allocated either (1) dynamically with `malloc`, `calloc`, etc. or (2) on the stack of a function, `bar` that either called `foo`, or calls a function that calls `foo`, or calls a function that calls a function that calls `foo`...

3. **Interpret your fields to perform GET, PUT logic or return a response indicating the invalidity in the request.** Finally, think about how to implement the logic in response to a client request. *Think: How can you leverage our helper functions to help you with this task?*

Implementation Approach

Below, we describe steps that can be used to implement your server.

Step 1: Build an Echo Server We suggest that you start building your server by building the Scaffolding. The scaffolding can work by simply repeating the bytes that the client writes to it. This ensures that your code is able to listen, accept, read from, and close connections.

Step 2: Sending Responses Next, we suggest that you design code that allows your server to send a response. This is important because it will be impossible to test whether your system can parse requests without being sure that your server can produce responses. Our suggestion is to unit test this: use the `cse130_nc` script to validate that your server produces correct responses before you move on to full-scale testing.

Step 3: Request Logic Then, we suggest that you start building code that performs **GET** and **PUT** requests. You might look back at your **memory** assignment and use our helper functions. At this point, you should be able to use any of the tests in Resources that handle "happy-path" scenarios. In other words: your code should work for tests that use requests that are supposed to return a **Status-Code** of 200 or 201.

Step 4: Parsing Requests Finally, we suggest that you start building code parses requests. This is the most complex task in this assignment and the one that trips up the most students. After this step, your server should be able to support all of the tests in Resources.